

# Assignment-1 Report

Astitva Aryan (22110048), Dhruv Sharma (22110074)

September 15, 2025

**Source Code Repository:** [github.com/astitvaaryan/CS-331...](https://github.com/astitvaaryan/CS-331...)

## Task 1

The purpose of this task was to implement a custom DNS resolution system consisting of a client and a server. The client was designed to parse DNS query packets from a given PCAP file  $[(041 + 074) \% 10 = 5]$  ('5.pcap'). For each query, the client generated a custom 8-byte header in the format HHMMSSID (Hour, Minute, Second, Sequential ID) and prepended it to the original DNS packet.

The modified packet was then sent to a custom server. The server, instead of performing a live DNS lookup, implemented a set of predefined, time-based rules. It used the hour and ID from the custom header to deterministically assign an IP address from a static pool of 15 addresses. This process simulates a time-based load balancing system.

## Methodology

The implementation consists of two Python scripts: a client and a server. The client is responsible for packet parsing and forwarding, while the server handles the custom logic for IP address assignment.

### Client Code (client.py)

```
1 import socket
2 from datetime import datetime
3 from scapy.all import rdpcap, DNS, DNSQR, raw
4
5 # --- Configuration ---
6 PCAP_FILE = '5.pcap'
7 SERVER_IP = '127.0.0.1'
8 SERVER_PORT = 9999
9 REPORT_FILE = 'report.txt'
10
11 def process_dns_queries():
12     """
13     Reads a pcap file, filters DNS queries, sends them to the custom server,
14     receives the resolved IP, and generates a final report.
15     """
16     print(f"[*] Reading packets from '{PCAP_FILE}'...")
17     try:
18         packets = rdpcap(PCAP_FILE)
19     except FileNotFoundError:
20         print(f"[ERROR] The file '{PCAP_FILE}' was not found. Please place it
    in the same directory.")
```

```

21         return
22
23     # Filter for packets that are standard DNS queries (qr=0 means query)
24     dns_queries = [pkt for pkt in packets if pkt.haslayer(DNS) and pkt[DNS].qr
25                     == 0]
26
27     if not dns_queries:
28         print("[*] No DNS queries found in the pcap file.")
29         return
30
31     print(f"[*] Found {len(dns_queries)} DNS queries. Processing...")
32
33     report_data = []
34     query_id_counter = 0
35
36     for query_packet in dns_queries:
37         try:
38             # 1. Create the custom 8-byte header (HHMMSSID).
39             now = datetime.now()
40             hhmss = now.strftime("%H%M%S")
41             # Format ID with a leading zero if it's a single digit (e.g., 00,
42             # 01, ..., 09, 10)
43             query_id_str = f"{query_id_counter:02d}"
44             custom_header = f"{hhmss}{query_id_str}"
45
46             # 2. Construct the payload: custom header + original DNS packet
47             bytes.
48             header_bytes = custom_header.encode('utf-8')
49             packet_bytes = raw(query_packet)
50             payload = header_bytes + packet_bytes
51
52             # 3. Connect to the server and send the payload.
53             with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
54                 s.connect((SERVER_IP, SERVER_PORT))
55                 s.sendall(payload)
56
57             # 4. Receive the server's response (the assigned IP address).
58             response = s.recv(1024).decode('utf-8')
59             assigned_ip = response
60
61             # 5. Log data for the report.
62             queried_domain = query_packet[DNSQR].qname.decode('utf-8').rstrip('.')
63
64             report_data.append((custom_header, queried_domain, assigned_ip))
65             print(f" - Sent Query for '{queried_domain}' with Header '{custom_header}' -> Got IP:{assigned_ip}")
66
67             query_id_counter += 1
68
69         except Exception as e:
70             print(f"[ERROR] Failed to process a query: {e}")
71
72     # 6. Generate the final report file.
73     generate_report(report_data)
74
75 def generate_report(data):
76     """Writes the final report to a text file and prints it."""
77     header = f"'Custom Header (HHMMSSID)':<30} | {'Domain Name':<40} | {'Assigned IP Address'}\n"
78     separator = "-" * 100 + "\n"
79
80     report_content = header + separator

```

```

78     for entry in data:
79         report_content += f"{entry[0]:<30} | {entry[1]:<40} | {entry[2]}\n"
80
81     try:
82         with open(REPORT_FILE, 'w') as f:
83             f.write(report_content)
84             print(f"\n[*] Report successfully generated: '{REPORT_FILE}')"
85     except IOError as e:
86         print(f"\n[ERROR] Could not write report to file: {e}")
87
88     print("\n--- Final Report ---")
89     print(report_content)
90
91
92 if __name__ == "__main__":
93     process_dns_queries()

```

Listing 1: The client script reads DNS queries from a PCAP file, adds a custom header, and sends them to the server for resolution.

## Server Code (server.py)

```

1  import socket
2  import threading
3
4  # Define the static IP pool as specified in the rules document.
5  IP_POOL = [
6      "192.168.1.1", "192.168.1.2", "192.168.1.3", "192.168.1.4", "192.168.1.5",
7      "192.168.1.6", "192.168.1.7", "192.168.1.8", "192.168.1.9", "192.168.1.10",
8      "192.168.1.11", "192.168.1.12", "192.168.1.13", "192.168.1.14", "
9      192.168.1.15"
10 ]
11
12 def handle_client(conn, addr):
13     """
14     Handles an individual client connection, applies the DNS resolution rules,
15     and sends back the assigned IP address.
16     """
17     print(f"[NEW CONNECTION] {addr} connected.")
18     try:
19         # Receive the payload from the client
20         data = conn.recv(2048)
21         if not data:
22             return
23
24         # 1. Extract the 8-byte custom header from the payload.
25         custom_header = data[:8].decode('utf-8')
26
27         # 2. Parse the header to get the hour (HH) and ID.
28         hour = int(custom_header[0:2])
29         query_id = int(custom_header[6:8])
30
31         # 3. Apply the time-based routing rules to find the starting index.
32         ip_pool_start = 0
33         if 4 <= hour <= 11: # Morning: 04:00-11:59
34             ip_pool_start = 0
35             time_of_day = "Morning"
36         elif 12 <= hour <= 19: # Afternoon: 12:00-19:59
37             ip_pool_start = 5
38             time_of_day = "Afternoon"
39         else: # Night: 20:00-03:59
40             ip_pool_start = 10

```

```

40         time_of_day = "Night"
41
42         # 4. Calculate the offset using the ID modulo 5.
43         offset = query_id % 5
44
45         # 5. Calculate the final index and select the IP address.
46         final_index = ip_pool_start + offset
47         assigned_ip = IP_POOL[final_index]
48
49         print(f" [PROCESSING] Header: {custom_header} | Hour: {hour} ({
time_of_day}) | ID: {query_id}")
50         print(f" [LOGIC] Pool Start: {ip_pool_start} | Offset (ID % 5): {
offset} | Final Index: {final_index}")
51         print(f" [RESULT] Assigned IP: {assigned_ip}")
52
53         # 6. Send the assigned IP back to the client.
54         conn.sendall(assigned_ip.encode('utf-8'))
55
56     except (ValueError, IndexError) as e:
57         print(f"[ERROR] Invalid data received from {addr}: {e}")
58         error_message = "Error: Invalid header or data format"
59         conn.sendall(error_message.encode('utf-8'))
60     finally:
61         conn.close()
62         print(f"[CONNECTION CLOSED] {addr}")
63
64 def start_server(host='127.0.0.1', port=9999):
65     """
66     Starts the DNS resolver server, listening for incoming connections.
67     """
68     server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
69     server_socket.bind((host, port))
70     server_socket.listen()
71     print(f"[LISTENING] Server is listening on {host}:{port}")
72
73     while True:
74         conn, addr = server_socket.accept()
75         # Create a new thread to handle each client connection so the server
can
76         # handle multiple requests concurrently without blocking.
77         thread = threading.Thread(target=handle_client, args=(conn, addr))
78         thread.start()
79
80 if __name__ == "__main__":
81     start_server()

```

Listing 2: The server script listens for client connections, parses the custom header, and assigns an IP based on predefined time-based rules.

## Output

```
dhruvs@dhruv-tufgamingfx504gdfx80gd in /home/dhruv/Semester 7/Networks/Assignment 1 via v3.13.7 (.venv) took 3s933ms
> python3 client.py
[*] Reading packets from '5.pcap'...
[*] Found 23 DNS queries. Processing...
- Sent Query for 'apple.com' with Header '01160200' -> Got IP: 192.168.1.11
- Sent Query for '_apple-mobdev._tcp.local' with Header '01160201' -> Got IP: 192.168.1.12
- Sent Query for '_apple-mobdev._tcp.local' with Header '01160202' -> Got IP: 192.168.1.13
- Sent Query for 'facebook.com' with Header '01160203' -> Got IP: 192.168.1.14
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160204' -> Got IP: 192.168.1.15
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160205' -> Got IP: 192.168.1.11
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160206' -> Got IP: 192.168.1.12
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160207' -> Got IP: 192.168.1.13
- Sent Query for 'amazon.com' with Header '01160208' -> Got IP: 192.168.1.14
- Sent Query for '_apple-mobdev._tcp.local' with Header '01160209' -> Got IP: 192.168.1.15
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160210' -> Got IP: 192.168.1.11
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160211' -> Got IP: 192.168.1.12
- Sent Query for 'twitter.com' with Header '01160212' -> Got IP: 192.168.1.13
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160213' -> Got IP: 192.168.1.14
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160214' -> Got IP: 192.168.1.15
- Sent Query for '_apple-mobdev._tcp.local' with Header '01160215' -> Got IP: 192.168.1.11
- Sent Query for '_apple-mobdev._tcp.local' with Header '01160216' -> Got IP: 192.168.1.12
- Sent Query for 'wikipedia.org' with Header '01160217' -> Got IP: 192.168.1.13
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160218' -> Got IP: 192.168.1.14
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160219' -> Got IP: 192.168.1.15
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160220' -> Got IP: 192.168.1.11
- Sent Query for 'Brother MFC-7860DW._pdl-datastream._tcp.local' with Header '01160221' -> Got IP: 192.168.1.12
- Sent Query for 'stackoverflow.com' with Header '01160222' -> Got IP: 192.168.1.13
```

Figure 1: Client Output showing the script processing DNS queries and receiving IPs.

## Results

The following table presents the output generated by the client after processing all DNS queries from the PCAP file and receiving the assigned IP addresses from the server.

Custom Header (HHMMSSID)	Domain Name	Assigned IP Add
02110100	apple.com	192.168.1.11
02110101	_apple-mobdev._tcp.local	192.168.1.12
02110102	_apple-mobdev._tcp.local	192.168.1.13
02110103	facebook.com	192.168.1.14
02110104	Brother	192.168.1.15
	MFC-7860DW._pdl-datastream._tcp.local	
02110105	Brother	192.168.1.11
	MFC-7860DW._pdl-datastream._tcp.local	
02110106	Brother	192.168.1.12
	MFC-7860DW._pdl-datastream._tcp.local	
02110107	Brother	192.168.1.13
	MFC-7860DW._pdl-datastream._tcp.local	
02110108	amazon.com	192.168.1.14
02110109	_apple-mobdev._tcp.local	192.168.1.15
02110110	Brother	192.168.1.11
	MFC-7860DW._pdl-datastream._tcp.local	
02110111	Brother	192.168.1.12
	MFC-7860DW._pdl-datastream._tcp.local	
02110112	twitter.com	192.168.1.13

Continued on next p

Table 1 – continued from previous page

Custom Header (HHMMSSID)	Domain Name	Assigned IP Addr
02110113	Brother	192.168.1.14
02110114	MFC-7860DW._pdl-datastream._tcp.local	192.168.1.15
	Brother	
02110115	MFC-7860DW._pdl-datastream._tcp.local	192.168.1.11
	_apple-mobdev._tcp.local	
02110116	_apple-mobdev._tcp.local	192.168.1.12
02110117	wikipedia.org	192.168.1.13
02110118	Brother	192.168.1.14
	MFC-7860DW._pdl-datastream._tcp.local	
02110119	Brother	192.168.1.15
	MFC-7860DW._pdl-datastream._tcp.local	
02110120	Brother	192.168.1.11
	MFC-7860DW._pdl-datastream._tcp.local	
02110121	Brother	192.168.1.12
	MFC-7860DW._pdl-datastream._tcp.local	
02110122	stackoverflow.com	192.168.1.13

## Task-2: Traceroute Protocol Behavior

The purpose of this task is to understand how the traceroute utility works in different operating systems by capturing and analyzing the network traffic it generates.

### Experiment on Windows

The `tracert` command was executed on Windows to trace the route to `www.google.com`. The network traffic was captured simultaneously using Wireshark.

```
C:\Users\Astitva_Aryan_>tracert www.google.com

Tracing route to www.google.com [142.251.42.68]
over a maximum of 30 hops:

  1    5 ms    5 ms    7 ms  10.7.0.5
  2    6 ms    5 ms    4 ms  172.16.4.7
  3   13 ms   11 ms   17 ms  14.139.98.1
  4    6 ms    4 ms    5 ms  10.117.81.253
  5   36 ms   15 ms   15 ms  10.154.8.137
  6   15 ms   16 ms   14 ms  10.255.239.170
  7   13 ms   12 ms   11 ms  10.152.7.214
  8   12 ms   13 ms   12 ms  72.14.204.62
  9   13 ms   15 ms   14 ms  72.14.239.103
 10   15 ms   14 ms   14 ms  142.251.69.105
 11   20 ms   16 ms   18 ms  bom12s21-in-f4.1e100.net [142.251.42.68]

Trace complete.
```

Figure 2: Output of `tracert www.google.com` on Windows.

## Experiment on Linux

The purpose of this experiment was to analyze the behavior of the ‘traceroute’ utility on a Linux-based operating system (Arch Linux). The process involved running a traceroute to a public destination (‘google.com’) while simultaneously capturing the generated network traffic using ‘tcpdump’. The captured data is used to answer fundamental questions about the tool’s underlying protocol mechanics.

### Methodology: Execution and Capture

The analysis was performed by running ‘traceroute google.com’ in one terminal and ‘sudo tcpdump -i wlo1 -n -vv ’host google.com or icmp’ in a second terminal. This allowed for a live capture of both the outgoing probes sent by the client and the incoming ICMP responses from routers along the path.

### Traceroute Output

The following is the terminal output from the ‘traceroute’ command. It successfully maps the path to the destination, showing a failure at hop 9.

```
1 traceroute to google.com (142.251.222.78), 30 hops max, 60 byte packets
2 1  10.1.144.3 (10.1.144.3)  2.815 ms  2.792 ms  2.781 ms
3 2  172.16.4.7 (172.16.4.7)  2.062 ms  2.756 ms  2.744 ms
4 3  14.139.98.1 (14.139.98.1)  3.861 ms  5.453 ms  5.435 ms
5 4  10.117.81.253 (10.117.81.253)  2.667 ms  2.620 ms  2.596 ms
6 5  10.154.8.137 (10.154.8.137)  11.950 ms  13.502 ms  12.397 ms
7 6  10.255.239.170 (10.255.239.170)  15.985 ms  14.107 ms  13.614 ms
8 7  10.152.7.214 (10.152.7.214)  9.616 ms  9.607 ms  9.597 ms
9 8  72.14.204.62 (72.14.204.62)  14.436 ms  14.422 ms  *
10 9  * * *
11 10  192.178.86.200 (192.178.86.200)  16.134 ms  142.251.77.100 (142.251.77.100)
    38.084 ms  38.075 ms
12 11  192.178.110.208 (192.178.110.208)  10.357 ms  192.178.110.104
    (192.178.110.104)  11.596 ms  192.178.110.110 (192.178.110.110)  38.046 ms
13 12  142.250.208.227 (142.250.208.227)  14.906 ms  pnbomb-bp-in-f14.1e100.net
    (142.251.222.78)  10.806 ms  142.250.209.71 (142.250.209.71)  12.880 ms
```

Listing 3: Terminal output of the traceroute command to google.com.

### Tcpdump Packet Capture

The following is an abridged ‘tcpdump’ capture that highlights the key packets involved in the traceroute process.

```
1 tcpdump: listening on wlo1, link-type EN10MB (Ethernet), snapshot length 262144
  bytes
2
3 % --- Probe for Hop 1 and its reply ---
4 01:35:40.916080 IP (tos 0x0, ttl 255, id 61206, proto ICMP (1), length 56)
5     10.1.144.3 > 10.1.144.19: ICMP time exceeded in-transit, length 36
6         IP (tos 0x0, ttl 1, id 61206, proto UDP (17), length 60)
7             10.1.144.19.60238 > 142.251.222.78.33434: UDP, length 32
8 ...
9 % --- Probe for Hop 2 and its reply ---
10 01:35:40.929641 IP (tos 0x0, ttl 248, id 63877, proto ICMP (1), length 168)
11     10.255.239.170 > 10.1.144.19: ICMP time exceeded in-transit, length 148
12         IP (tos 0x0, ttl 2, id 42881, proto UDP (17), length 60)
13             10.1.144.19.38871 > 142.251.222.78.33449: [udp sum ok] UDP, length
14         32
    ...
```

```

15 % --- Reply from the Final Destination ---
16 01:35:41.053762 IP (tos 0x80, ttl 115, id 0, proto ICMP (1), length 56)
17   142.251.222.78 > 10.1.144.19: ICMP 142.251.222.78 udp port 33474
   unreachable, length 36
18   IP (tos 0x80, ttl 3, id 48364, proto UDP (17), length 60)
19   10.1.144.19.49232 > 142.251.222.78.33474: UDP, length 32
20
21 39 packets captured

```

Listing 4: Tcpdump capture showing UDP probes and ICMP replies.

## Analysis and Questions

The following questions are answered based on the captured traffic and command-line output.

1. What protocol does Windows `tracert` use by default, and what protocol does Linux `traceroute` use by default?

Windows `tracert` uses the **ICMP** (Internet Control Message Protocol) by default. It sends a series of ICMP Echo Request packets with increasing TTL values.

Linux `traceroute` uses **UDP** (User Datagram Protocol) by default. It sends UDP packets to high-numbered ports.

No.	Time	Source	Destination	Protocol	Length	Info
23472	934.737220	172.16.4.7	10.7.8.129	ICMP	120	Time-to-live exceeded (Time to live exceeded in transit)
23473	934.778648	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23474	934.819482	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23475	934.819482	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23476	934.820755	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23477	934.820755	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23478	934.898433	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23479	934.899213	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23480	934.899213	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23481	934.900063	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23482	934.980461	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23483	934.985273	10.7.8.129	13.89.179.11	TCP	66	49947 → 443 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=256 SACK
23484	934.985720	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23485	934.985720	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23486	934.985720	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23487	935.076176	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23488	935.076176	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23489	935.076176	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23490	935.076176	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23491	935.090371	10.7.8.129	155.133.225.20	TLSv1.2	108	Application Data
23492	935.141610	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23493	935.141610	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23494	935.142950	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23495	935.142950	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23496	935.178032	155.133.225.20	10.7.8.129	TCP	54	27020 → 53174 [ACK] Seq=1358 Ack=2037 Win=8195 Len=0
23497	935.219262	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)
23498	935.220091	10.7.0.5	10.7.8.129	ICMP	70	Redirect (Redirect for host)

Figure 3: Wireshark capture filtered for ICMP traffic, showing the Echo Requests and Time-to-live exceeded messages generated by `tracert`.

2. Some hops in your `traceroute` output may show `***`. Provide at least two reasons why a router might not reply.

The asterisks (`***`) indicate that no reply was received from a router for a given probe. Two primary reasons for this are:

- **Firewall Configuration:** The router, or a firewall protecting it, is explicitly configured to drop ICMP packets and not send any ICMP error messages. This is a common security practice to prevent network mapping.
- **ICMP Rate Limiting:** To prevent denial-of-service attacks or reduce load, routers are often configured to limit the rate at which they send ICMP error messages. If probes arrive during a period where this limit is exceeded, the router will simply ignore them.



```

dhruvs@dhruv-tufgamingfx504gdfx80gd in ~ as 🐼 took 14s649ms
# ) sudo tcpdump -i wlo1 -n -vv 'host google.com or icmp'
tcpdump: listening on wlo1, link-type EN10MB (Ethernet), snapshot length 262144 bytes
01:35:40.915374 IP (tos 0xc0, ttl 254, id 60351, offset 0, flags [none], proto ICMP (1), length 88)
  172.16.4.7 > 10.1.144.19: ICMP time exceeded in-transit, length 68
    IP (tos 0x0, ttl 1, id 49334, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.49377 > 142.251.222.78.33437: [udp sum ok] UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 255, id 61206, offset 0, flags [none], proto ICMP (1), length 56)
  10.1.144.3 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 61206, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.60238 > 142.251.222.78.33434: UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 254, id 60352, offset 0, flags [none], proto ICMP (1), length 88)
  172.16.4.7 > 10.1.144.19: ICMP time exceeded in-transit, length 68
    IP (tos 0x0, ttl 1, id 25452, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.49775 > 142.251.222.78.33438: [udp sum ok] UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 254, id 60353, offset 0, flags [none], proto ICMP (1), length 88)
  172.16.4.7 > 10.1.144.19: ICMP time exceeded in-transit, length 68
    IP (tos 0x0, ttl 1, id 35697, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.49009 > 142.251.222.78.33439: [udp sum ok] UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 255, id 63786, offset 0, flags [none], proto ICMP (1), length 56)
  10.1.144.3 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 63786, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.47396 > 142.251.222.78.33435: UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 255, id 61846, offset 0, flags [none], proto ICMP (1), length 56)
  10.1.144.3 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 61846, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.58246 > 142.251.222.78.33436: UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 252, id 0, offset 0, flags [none], proto ICMP (1), length 56)
  10.117.81.253 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 41392, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.50512 > 142.251.222.78.33444: UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 252, id 0, offset 0, flags [none], proto ICMP (1), length 56)
  10.117.81.253 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 48979, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.47392 > 142.251.222.78.33443: UDP, length 32
01:35:40.916080 IP (tos 0xc0, ttl 252, id 0, offset 0, flags [none], proto ICMP (1), length 56)
  10.117.81.253 > 10.1.144.19: ICMP time exceeded in-transit, length 36
    IP (tos 0x0, ttl 1, id 31218, offset 0, flags [none], proto UDP (17), length 60)
  10.1.144.19.38192 > 142.251.222.78.33445: UDP, length 32

```

Figure 4: Linux TCPDUMP capture showing the Echo Requests and Time-to-live exceeded messages generated by traceroute.

### 3. In Linux traceroute, which field in the probe packets changes between successive probes sent to the destination?

#### Linux Analysis

The analysis was performed by running ‘traceroute google.com’ in one terminal and ‘sudo tcpdump -i wlo1 -n -vv ‘host google.com or icmp’ in a second terminal. This allowed for a live capture of both the outgoing probes and the incoming responses.

#### Outputs

##### Traceroute Output

```

1 traceroute to google.com (142.251.222.78), 30 hops max, 60 byte packets
2 1 10.1.144.3 (10.1.144.3) 2.815 ms 2.792 ms 2.781 ms
3 2 172.16.4.7 (172.16.4.7) 2.062 ms 2.756 ms 2.744 ms
4 3 14.139.98.1 (14.139.98.1) 3.861 ms 5.453 ms 5.435 ms
5 4 10.117.81.253 (10.117.81.253) 2.667 ms 2.620 ms 2.596 ms
6 5 10.154.8.137 (10.154.8.137) 11.950 ms 13.502 ms 12.397 ms
7 6 10.255.239.170 (10.255.239.170) 15.985 ms 14.107 ms 13.614 ms
8 7 10.152.7.214 (10.152.7.214) 9.616 ms 9.607 ms 9.597 ms
9 8 72.14.204.62 (72.14.204.62) 14.436 ms 14.422 ms *
10 9 * * *
11 10 192.178.86.200 (192.178.86.200) 16.134 ms 142.251.77.100
    (142.251.77.100) 38.084 ms 38.075 ms
12 11 192.178.110.208 (192.178.110.208) 10.357 ms 192.178.110.104
    (192.178.110.104) 11.596 ms 192.178.110.110 (192.178.110.110) 38.046
    ms

```

```

13 12 142.250.208.227 (142.250.208.227) 14.906 ms pnbomb-bp-in-f14.1e100.net
    (142.251.222.78) 10.806 ms 142.250.209.71 (142.250.209.71) 12.880 ms
14

```

Listing 5: Terminal output of the traceroute command to google.com.

## Tcpdump Packet Capture

```

1 tcpdump: listening on wlo1, link-type EN10MB (Ethernet), snapshot length
  262144 bytes
2 01:35:40.916080 IP (tos 0x0, ttl 255, id 61206, offset 0, flags [none],
  proto ICMP (1), length 56)
3   10.1.144.3 > 10.1.144.19: ICMP time exceeded in-transit, length 36
4   IP (tos 0x0, ttl 1, id 61206, offset 0, flags [none], proto UDP
  (17), length 60)
5       10.1.144.19.60238 > 142.251.222.78.33434: UDP, length 32
6   ...
7 01:35:40.929641 IP (tos 0x0, ttl 248, id 63877, offset 0, flags [none],
  proto ICMP (1), length 168)
8   10.255.239.170 > 10.1.144.19: ICMP time exceeded in-transit, length 148
9   IP (tos 0x0, ttl 2, id 42881, offset 0, flags [none], proto UDP
  (17), length 60)
10      10.1.144.19.38871 > 142.251.222.78.33449: [udp sum ok] UDP,
  length 32
11   ...
12 01:35:41.053762 IP (tos 0x80, ttl 115, id 0, offset 0, flags [none], proto
  ICMP (1), length 56)
13   142.251.222.78 > 10.1.144.19: ICMP 142.251.222.78 udp port 33474
  unreachable, length 36
14   IP (tos 0x80, ttl 3, id 48364, offset 0, flags [none], proto UDP
  (17), length 60)
15      10.1.144.19.49232 > 142.251.222.78.33474: UDP, length 32
16   ...
17 39 packets captured
18

```

Listing 6: Abridged tcpdump capture showing UDP probes and ICMP replies.

**Analysis and Findings** The captured data shows that two fields change for each successive probe:

- **IP Time To Live (TTL):** The TTL field in the IP header is incremented by one for each hop ('ttl 1', 'ttl 2', 'ttl 3', etc.). This is the core mechanism of traceroute.
- **UDP Destination Port:** The destination port in the UDP header is also incremented for each probe ('...33434', '...33435', etc.). This helps to distinguish between replies and match them to the correct outgoing probe.

#### 4. At the final hop, how is the response different compared to the intermediate hop?

- **Intermediate Hops (Both Windows & Linux):** Reply with an ICMP message of Type 11, Code 0, which is **"Time-to-live exceeded."** This is an error message indicating the packet expired before reaching the destination.
- **Final Hop (Windows tracert):** Since the probe is an ICMP Echo Request, the destination replies with a successful ICMP message of Type 0, Code 0, which is an **"Echo Reply."**

1150	70.964138	10.7.8.129	142.250.77.68	ICMP	106 Echo (ping) request	id=0x0001, seq=425/43265, ttl=11 (reply in 1152)
1152	70.977345	142.250.77.68	10.7.8.129	ICMP	106 Echo (ping) reply	id=0x0001, seq=425/43265, ttl=115 (request in 1150)
1153	70.981109	10.7.8.129	142.250.77.68	ICMP	106 Echo (ping) request	id=0x0001, seq=426/43521, ttl=11 (reply in 1155)
1155	71.072795	142.250.77.68	10.7.8.129	ICMP	106 Echo (ping) reply	id=0x0001, seq=426/43521, ttl=115 (request in 1153)

Figure 5: Windows Last Lines

- **Final Hop (Linux traceroute):** The UDP probe is sent to a high, unused port. Since no service is listening on that port, the destination’s OS sends back an ICMP message of Type 3, Code 3, which is **”Destination Unreachable (Port Unreachable).”** This error message signals to ‘traceroute’ that the final destination has been reached.

```

10.1.144.19.47891 > 142.251.222.78.33471: UDP, length 32
01:35:41.054185 IP (tos 0x00, ttl 115, id 0, offset 0, flags [none], proto ICMP (1), length 56)
142.251.222.78 > 10.1.144.19: ICMP 142.251.222.78 udp port 33468 unreachable, length 36
IP (tos 0x00, ttl 1, id 23247, offset 0, flags [none], proto UDP (17), length 60)
10.1.144.19.59675 > 142.251.222.78.33468: UDP, length 32
01:35:41.054185 IP (tos 0x00, ttl 245, id 11153, offset 0, flags [none], proto ICMP (1), length 96)
142.251.77.100 > 10.1.144.19: ICMP time exceeded in-transit, length 76
IP (tos 0x00, ttl 1, id 61318, offset 0, flags [none], proto UDP (17), length 60)
10.1.144.19.47458 > 142.251.222.78.33462: [udp sum ok] UDP, length 32
01:35:41.055721 IP (tos 0x00, ttl 115, id 0, offset 0, flags [none], proto ICMP (1), length 56)
142.251.222.78 > 10.1.144.19: ICMP 142.251.222.78 udp port 33476 unreachable, length 36
IP (tos 0x00, ttl 4, id 65005, offset 0, flags [none], proto UDP (17), length 60)
10.1.144.19.59392 > 142.251.222.78.33476: UDP, length 32
01:35:41.055722 IP (tos 0x00, ttl 115, id 0, offset 0, flags [none], proto ICMP (1), length 56)
142.251.222.78 > 10.1.144.19: ICMP 142.251.222.78 udp port 33472 unreachable, length 36
IP (tos 0x00, ttl 2, id 23394, offset 0, flags [none], proto UDP (17), length 60)
10.1.144.19.37766 > 142.251.222.78.33472: UDP, length 32

```

Figure 6: Linux Last Lines

5. **Suppose a firewall blocks UDP traffic but allows ICMP how would this affect the results of Linux traceroute vs. Windows tracert?**

A firewall blocking UDP but allowing ICMP would have different effects on the two systems:

- **Windows tracert:** Since Windows `tracert` uses ICMP for its probes, it would be **unaffected** by this firewall and would likely complete the trace successfully.
- **Linux traceroute:** The default Linux `traceroute` uses UDP packets for its probes. This trace would **fail** at the firewall. Its UDP probes would be blocked, and it would not receive any responses from hops beyond that point, resulting in a series of timeouts (‘\*\*\*’).

## References

1. Traceroute Explanation, available at: <https://share.google/4b7B4sdrua4iLotKd> (Accessed: 15 September 2025).